

FDE Client Discovery Questions

Client-Specific AI Copilot for Logistics Operations

A complete discovery checklist for Forward Deployment Engineers covering business problems, users, workflow, tracking, pricing, policies, escalation, AI decisions, integrations, security, RAG, monitoring, evaluation, deployment, governance, and production rollout.

246 client questions grouped into 20 practical discovery categories.

FDE mindset: Do not begin with 'Which LLM should we use?' First understand who has the problem, how the business works today, where the source of truth lives, what is safe to automate, what can fail, and how success will be measured.

1. Business problem and objective

1. What is the main business problem you want this Copilot to solve?
2. What are the biggest pain points in the current logistics support process?
3. Which tasks consume the most time for operations employees?
4. Which tasks create the most customer complaints?
5. Which tasks currently require switching between multiple applications?
6. What problems are you expecting AI to solve that your current systems cannot solve easily?
7. What would make this project successful from a business perspective?
8. Which business KPIs should improve after deployment?
9. Is the main objective faster response, lower cost, fewer errors, better customer experience, or all of these?
10. What is the expected ROI or business impact?

2. Users and personas

11. Who will use the Copilot?
12. Will it be used by operations staff, support agents, managers, customers, or all of them?
13. How many users do you expect initially?
14. How many users could use it after full rollout?
15. Are users located in one country or multiple regions?
16. What languages must the Copilot support?
17. Do different user roles need different capabilities?
18. Should managers see information that operators cannot see?
19. Do customers have different access rights from employees?
20. What level of technical knowledge do the users have?

3. Current workflow

21. Please walk me through the current process from receiving a logistics query to resolving it.
22. Which systems does an employee normally open to answer one customer question?
23. Which parts of this workflow are manual?
24. Which parts are already automated?
25. Where do employees usually get stuck?
26. What is the average handling time for a query today?
27. How are unusual or difficult cases handled?
28. At what point is an issue escalated to a supervisor?
29. Are there existing standard operating procedures for these workflows?
30. What information must be recorded after each interaction?

4. Shipment tracking

31. Which system is the source of truth for shipment information?
32. Are you using a TMS, WMS, ERP, or custom tracking system?

- 33. Does the tracking system expose a REST API?
- 34. If not, is database access available?
- 35. What identifier is used for shipments?
- 36. What formats can shipment IDs have?
- 37. What shipment statuses currently exist?
- 38. What fields should the Copilot return for a shipment?
- 39. Should users see the current location of the shipment?
- 40. Should estimated delivery time be displayed?
- 41. Should historical shipment events be shown?
- 42. Can one shipment contain multiple packages or containers?
- 43. Should users be able to search by customer ID, order ID, invoice number, or only shipment ID?
- 44. How frequently is shipment status updated?
- 45. Is real-time tracking required?

5. Delayed shipments and exceptions

- 46. How do you currently define a delayed shipment?
- 47. Is delay calculated automatically or manually?
- 48. Are there different thresholds depending on service type?
- 49. Should the Copilot automatically identify delayed shipments?
- 50. Should it explain the reason for a delay?
- 51. What delay reasons are currently captured?
- 52. Should high-value or critical shipments receive different handling?
- 53. Which types of exceptions require immediate escalation?
- 54. Should the Copilot proactively alert users about severe delays?
- 55. Who should receive such alerts?

6. Pricing and cost calculation

- 56. Where does your delivery-pricing logic currently live?
- 57. Is pricing calculated by ERP, TMS, an Excel sheet, or another pricing engine?
- 58. What variables influence the delivery cost?
- 59. Does cost depend on weight?
- 60. Does cost depend on distance?
- 61. Does volumetric weight apply?
- 62. Are there different service priorities such as standard, express, and same-day?
- 63. Are fuel surcharges applied?
- 64. Are taxes added separately?
- 65. Do different customers have negotiated rates?
- 66. Are discounts applied based on volume or contract?
- 67. Are there minimum charges?
- 68. Are there different prices for different routes or regions?

- 69. Should the Copilot calculate the final price itself or call the existing pricing system?
- 70. Who is responsible if a price shown by the Copilot is incorrect?

7. Policies and knowledge base

- 71. Where are logistics policies currently stored?
- 72. Are they in PDF, Word, SharePoint, Confluence, email, or another knowledge base?
- 73. Which document repository is considered authoritative?
- 74. How often do policies change?
- 75. Who approves policy changes?
- 76. Are there multiple versions of the same policy?
- 77. Should the Copilot only use the latest approved version?
- 78. Should answers include citations to the original policy?
- 79. Should employees be able to open the source document from the answer?
- 80. Are some policies confidential?
- 81. Are some policies only visible to managers?
- 82. What should the Copilot do when two policies conflict?
- 83. What should happen if no policy supports the answer?
- 84. Should the system refuse to answer when supporting evidence cannot be found?

8. Escalation and human handoff

- 85. Which system is currently used for escalations?
- 86. Is it ServiceNow, Zendesk, Jira, email, or another helpdesk?
- 87. What conditions should automatically create an escalation?
- 88. Should users also be able to explicitly request escalation?
- 89. What information must be included in the escalation ticket?
- 90. Should the conversation history be attached?
- 91. Should shipment details automatically be added?
- 92. How should priority be determined?
- 93. Which team should receive each type of escalation?
- 94. What SLA exists for escalated cases?
- 95. Should the Copilot display the ticket number to the user?
- 96. How should duplicate escalation tickets be prevented?
- 97. Should the user be able to check the escalation status later?

9. Agent and AI decision making

- 98. Which user intents should the Copilot support initially?
- 99. Which queries should always be answered by deterministic rules?
- 100. Which queries could use an LLM?
- 101. Should the Copilot automatically select tools?
- 102. Are there actions that should always require human approval?

- 103. Should the agent be allowed to call more than one tool for a single question?
- 104. What should happen when the user's intent is unclear?
- 105. What confidence threshold should trigger escalation?
- 106. Should the system ask a clarifying question when information is missing?
- 107. Which actions must never be executed automatically?
- 108. Should the user be shown which tool was used?
- 109. Do you need explainability for agent decisions?
- 110. Should every agent decision be auditable?

10. Integration requirements

- 111. Which existing enterprise systems must be integrated?
- 112. Which APIs are already available?
- 113. Do you have API documentation?
- 114. What authentication method do those APIs use?
- 115. Are there API rate limits?
- 116. What is the expected API response time?
- 117. What happens when an external API is unavailable?
- 118. Are sandbox/test environments available?
- 119. Can we use test credentials?
- 120. Is there an API gateway already in the organization?
- 121. Are any systems accessible only inside the corporate network?
- 122. Will VPN or private connectivity be required?
- 123. Are there firewalls or IP allow-lists?
- 124. Do APIs have versioning policies?

11. Authentication and authorization

- 125. How do employees authenticate today?
- 126. Do you use Microsoft Entra ID, Okta, Keycloak, or another identity provider?
- 127. Is OAuth2/OIDC supported?
- 128. Should the Copilot use enterprise SSO?
- 129. What user roles exist?
- 130. What permissions should each role have?
- 131. Should users only be able to see shipments belonging to their region?
- 132. Are there customer-level access restrictions?
- 133. Should administrators have access to monitoring dashboards?
- 134. Do you need multi-factor authentication?
- 135. What should happen when an employee leaves the company?
- 136. How quickly should access be revoked?

12. Security and privacy

- 137. What types of sensitive information will pass through the Copilot?
- 138. Will the system process customer personal data?
- 139. Will it process addresses, phone numbers, or payment information?
- 140. Is there any information that must never be sent to an external LLM?
- 141. Are there data residency requirements?
- 142. Which security standards must the solution comply with?
- 143. Are GDPR, India DPDP Act, CCPA, PCI-DSS, or internal InfoSec policies applicable?
- 144. Should sensitive fields be masked or redacted?
- 145. How long should conversation logs be retained?
- 146. Who can view audit logs?
- 147. Should prompts and responses be stored?
- 148. Are users allowed to upload documents?
- 149. Should uploaded files be malware scanned?
- 150. What are your requirements for encryption at rest and in transit?

13. LLM and RAG questions

- 151. Are you comfortable using a cloud-hosted LLM?
- 152. Do you require a private or self-hosted model?
- 153. Which model providers are approved by security?
- 154. Can company data be sent outside the corporate network?
- 155. Should conversation data be used for model training?
- 156. Is retrieval-augmented generation required?
- 157. Which documents should be indexed?
- 158. How should documents be chunked and updated?
- 159. How quickly must new policies become searchable?
- 160. Should citations be mandatory?
- 161. What should happen when retrieval confidence is low?
- 162. Should the Copilot show the exact source used?
- 163. Do you require multilingual embeddings?

14. Reliability and resilience

- 164. What availability target do you expect?
- 165. Is 99.9% availability sufficient?
- 166. What is the acceptable response time?
- 167. What is the maximum acceptable latency?
- 168. How many concurrent users should the system support?
- 169. What are the peak traffic periods?
- 170. What should happen if the tracking system is unavailable?
- 171. Should cached shipment information be displayed?
- 172. How old can cached information be?

- 173. Should the application retry failed API calls?
- 174. How many retries are acceptable?
- 175. Should circuit breakers be implemented?
- 176. What is the expected disaster recovery strategy?
- 177. What recovery time objective, RTO, is required?
- 178. What recovery point objective, RPO, is required?

15. Monitoring and observability

- 179. Which monitoring platform does the organization currently use?
- 180. Do you use Grafana, Datadog, Splunk, Azure Monitor, CloudWatch, or another solution?
- 181. What technical metrics should we capture?
- 182. Should we monitor query count?
- 183. Should we monitor tool usage?
- 184. Should we monitor average latency and p95 latency?
- 185. Should we monitor agent-routing accuracy?
- 186. Should we track escalation rates?
- 187. What error rate should trigger an alert?
- 188. Who should receive alerts?
- 189. Should alerts go to email, Teams, Slack, PagerDuty, or another channel?
- 190. Do you need distributed tracing?
- 191. How long should logs be retained?
- 192. Do you need dashboards for business users as well as technical teams?

16. Business monitoring

- 193. What business KPIs should the Copilot improve?
- 194. What is the current average handling time?
- 195. What is the current first-contact resolution rate?
- 196. What is the current escalation rate?
- 197. What is current customer satisfaction or CSAT?
- 198. How much employee time do you expect the Copilot to save?
- 199. What automation rate would be considered successful?
- 200. What percentage of questions should be answered without human intervention?

17. Testing and evaluation

- 201. Do you have historical support queries we can use for testing?
- 202. Can we build a dataset of real operator questions?
- 203. Who will provide the expected correct answers?
- 204. How many test cases should be used before production?
- 205. What intent-routing accuracy is acceptable?
- 206. What tool-selection accuracy is acceptable?

- 207. What hallucination rate is acceptable?
- 208. How should we test unclear questions?
- 209. How should we test incorrect shipment IDs?
- 210. How should the system behave if required parameters are missing?
- 211. Should adversarial prompt-injection testing be included?
- 212. Should load testing be performed?
- 213. What maximum concurrent workload should be tested?

18. Deployment

- 214. Where should the application be deployed?
- 215. Is the preferred environment on-premises, AWS, Azure, GCP, or private cloud?
- 216. Is Kubernetes already available?
- 217. Can Docker containers be deployed?
- 218. Do you already have a container registry?
- 219. Is CI/CD already available?
- 220. Which CI/CD platform do you use?
- 221. Are separate development, staging, UAT, and production environments required?
- 222. Who approves production releases?
- 223. Are maintenance windows required?
- 224. Is blue-green or canary deployment preferred?
- 225. What rollback strategy is required?

19. Governance and ownership

- 226. Who owns the Copilot after deployment?
- 227. Who owns shipment integration?
- 228. Who owns pricing rules?
- 229. Who owns policy content?
- 230. Who owns escalation workflows?
- 231. Who approves AI-related changes?
- 232. Who reviews production incidents?
- 233. Who is responsible for updating policies?
- 234. Who is responsible for retraining or changing the agent?
- 235. What change-management process must be followed?

20. Rollout and adoption

- 236. Should we start with a small pilot?
- 237. Which team should participate in the pilot?
- 238. How many users should be included initially?
- 239. How long should the pilot run?
- 240. What criteria determine whether we expand the rollout?

- 241.** Do users need training?
- 242.** Should we provide user documentation?
- 243.** Should the system collect thumbs-up/thumbs-down feedback?
- 244.** Should employees be able to report incorrect answers?
- 245.** How will feedback be reviewed?
- 246.** Who will prioritize improvement requests?

Top 15 Questions for the First Client Meeting

Use these when you have only 30-45 minutes with the client. They help establish the problem, users, systems of record, integrations, security, success criteria, and production expectations.

1. What is the exact business problem we are trying to solve?
2. Who will use the Copilot?
3. What are the top 5 questions users ask today?
4. Walk me through the current workflow end to end.
5. Which systems do employees use today?
6. Which system is the source of truth for shipment tracking?
7. Where does pricing logic currently live?
8. Where are approved policies stored?
9. Where should unresolved issues be escalated?
10. Which APIs are available for these systems?
11. How do users authenticate today?
12. What sensitive data will pass through the system?
13. What response time and availability do you expect?
14. How will we measure whether the Copilot is successful?
15. What must be true before you would approve this solution for production?

Teaching takeaway: A strong FDE discovers the business first, designs around trusted systems of record, keeps automation safe, and treats monitoring and post-go-live improvement as part of the solution.